

The Emotional Spark - Concepts, Challenges, and Best Practices in Sentiment Analysis

Dr. Stefan Gindl, Freelancer

Data Scientist | Software Developer | Assistant Professor

stefan.gindl@gmail.com


- Concepts
- Applications
- Tools
- Features
- Evaluation & Limitations
- Corpus Annotation

Concepts

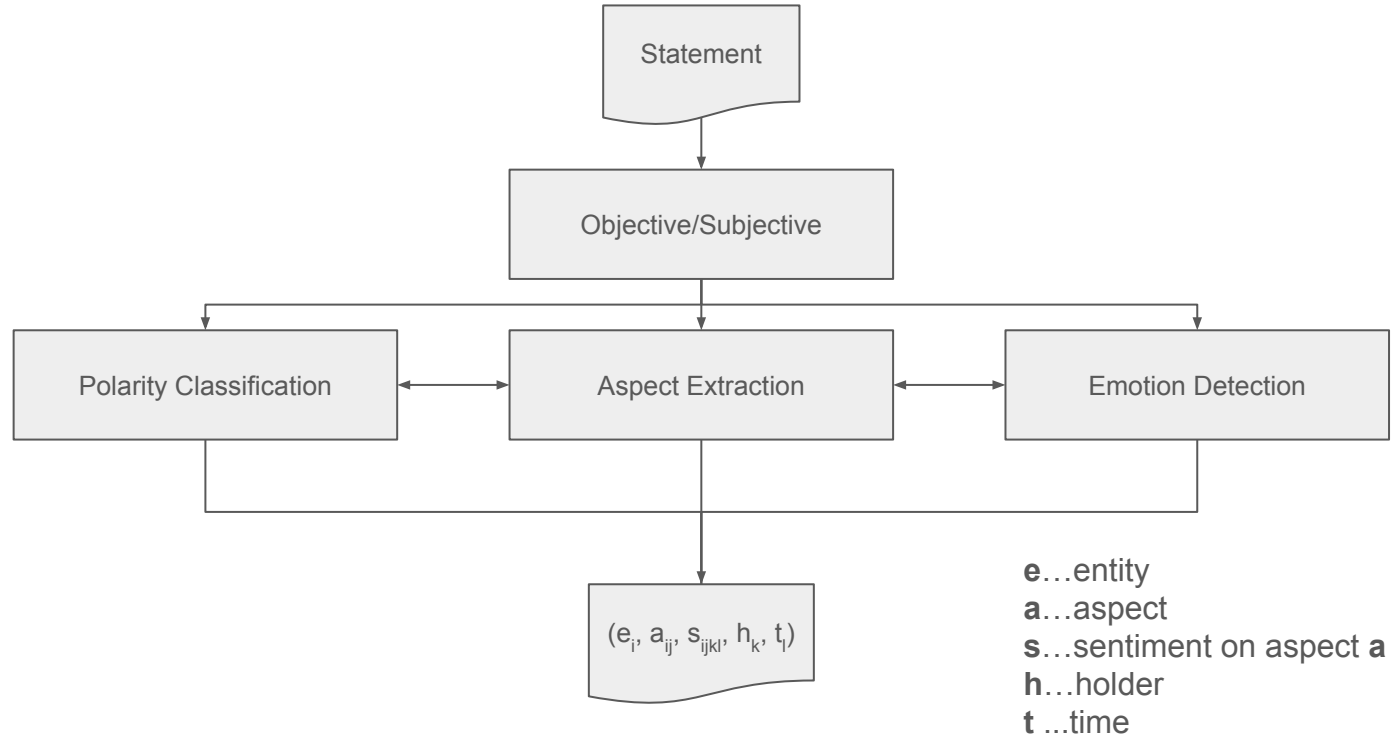
Concepts

- Sentiment analysis
 - = opinion mining
 - = sentiment detection
 - = polarity classification (?)
- Objective (factual) vs subjective (opinionated) statements
- Aspect-based opinion mining
- Emotion Detection
 - SentiSense
 - SenticNet

Robert Plutchik's wheel of emotions



Concepts



"The **iPhone** is very good, but they still need to work on **battery life** and **security** issues", John said yesterday.



entity	aspect	sentiment	holder	time
iPhone	GENERAL	pos	John	2019-06-18
iPhone	security	neg	John	2019-06-18
iPhone	battery life	neg	John	2019-06-18

Applications

Applications

- Media monitoring:
 - Get customer feedback from online media
 - Measure the political sentiment of the public
 - Health*:
 - Pharmaceutical companies mine online health information to monitor patients' opinions on their products and services, and to obtain feedback on their performance and the consumers' satisfaction.
 - Secondary effects of drugs, alternative therapies, undocumented symptoms
- As an additional feature for other applications, e.g. stock price prediction
- Built-in feature in CRM tools

* journals.plos.org/plosone/article?id=10.1371/journal.pone.0207996

Applications

Health Sector

- Algorithmic Detection and Analysis of **Vaccine-Denialist** Sentiment Clusters in **Social Networks** (2019): arxiv.org/abs/1905.12908
- Understanding Perceptions and Attitudes in **Breast Cancer Discussions on Twitter** (2019): arxiv.org/abs/1905.12469
- Distinguishing Clinical Sentiment: The Importance of Domain Adaptation in **Psychiatric Patient Health Records** (2019): arxiv.org/abs/1904.03225
- "Hang in There": Lexical and Visual Analysis to Identify Posts Warranting Empathetic Responses (2019): arxiv.org/abs/1903.05210
- Natural Language Processing, Sentiment Analysis and Clinical Analytics (2019): arxiv.org/abs/1902.00679

Applications

Finance

- **Market Trend Prediction** using Sentiment Analysis: Lessons Learned and Paths Forward (2019): arxiv.org/abs/1903.05440
- Predicting the Effects of **News Sentiments on the Stock Market** (2018): arxiv.org/abs/1812.04199
- Leveraging Financial News for **Stock Trend Prediction** with Attention-Based Recurrent Neural Network (2018): arxiv.org/abs/1811.06173
- Using Stock Prices as Ground Truth in Sentiment Analysis to Generate Profitable Trading Signals (2018): arxiv.org/abs/1811.02886

Applications

Politics

- Topic-Specific Sentiment Analysis Can Help **Identify Political Ideology** (2018): arxiv.org/abs/1810.12897
- Sentiment Aggregate Functions for **Political Opinion Polling** using Microblog Streams (2016): arxiv.org/abs/1606.05242
- Learning from the News: Predicting Entity Popularity on Twitter (2016): arxiv.org/abs/1607.03057
- Dynamic Allocation of Crowd Contributions for Sentiment Analysis during the 2016 U.S. Presidential Election (2016): arxiv.org/abs/1608.08953

Approaches

- Simplest: count number of positive and negative terms
- Feature engineering: filter irrelevant language particles, enrich useful particles with further information
- All sorts of machine learners
 - Naive Bayes, SVM, Linear Regression
- Deep Learning
 - Sentiment analysis solved?

Tools

Out-of-the-box

- Python
 - VADER: Valence Aware Dictionary and sEntiment Reasoner), a lexicon and rule-based sentiment analysis tool
 - NLTK: has a VADER wrapper
 - TextBlob:
 - spacy: AFAIK no
- Outside of Python
 - R: sentimentr (dictionary-based, uses valence shifters)
 - Java: Stanford CoreNLP

Out-of-the-box

```
[ ] 1 statement = ("This movie doesn't care about cleverness, wit or any other kind "  
2             "of intelligent humor.")
```

▼ NLTK

```
[ ] 1 from nltk.sentiment.vader import SentimentIntensityAnalyzer  
2 import nltk  
3 nltk.download('vader_lexicon')
```

```
[ ] 1 sa = SentimentIntensityAnalyzer()  
2 sa.polarity_scores(statement)
```

```
↳ {'compound': -0.1338, 'neg': 0.265, 'neu': 0.497, 'pos': 0.239}
```

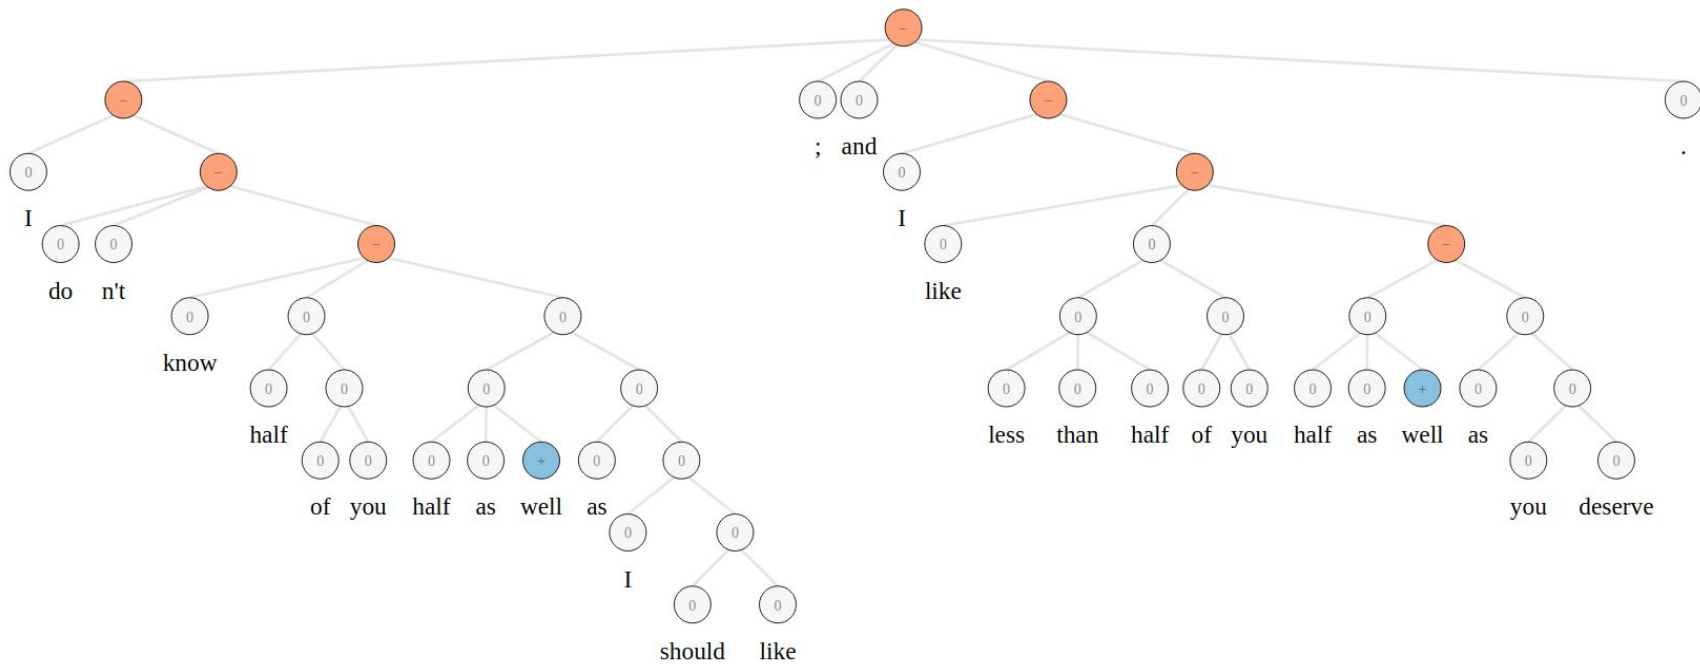
▼ TextBlob

```
[ ] 1 from textblob import TextBlob  
2 TextBlob(statement).polarity
```

```
↳ 0.425
```

Let's try it out...

→ nlp.stanford.edu/sentiment/



Fun with Sentiment Analysis

→ colab.research.google.com/drive/1umsBCj97qn07Sb76YO1N3H9ADnTFj8w_

Features

Features

- Bag-of-Words, e.g. counts or TF-IDF
- Stopword filtering
- Sentiment lexicons
 - General Inquirer, Opinion Lexicon, SentiSense and SentiStrength, SentiWordNet, SenticNet


Stopwords

```
[ ] 1 from nltk.corpus import stopwords
    2 import nltk
    3 import pandas as pd
    4 import requests
    5
    6 nltk.download('stopwords')
    7 github_stopwords_en = ('https://raw.githubusercontent.com/stopwords-iso/'
    8                        'stopwords-en/master/stopwords-en.txt')
    9
   10 stopwords_iso = requests.get(github_stopwords_en).text.split('\n')
   11 stopwords_nltk = set(stopwords.words("english"))
```

```
[ ] 1 print(f'Examples of stopwords_iso: {stopwords_iso[:5]}')
    2 print(f'Examples of stopwords_nltk: {list(stopwords_nltk)[:5]}')
```

➞ Examples of stopwords_iso: ["'ll", "'tis", "'twas", "'ve", '10']
Examples of stopwords_nltk: ['wouldn', 'the', 'and', 'by', 'she']

```
[ ] 1 stopwords = pd.Series({'stopwords_iso': len(stopwords_iso),
    2                      'stopwords_nltk': len(stopwords_nltk)})
    3 stopwords.plot.bar()
```



Sentiment Lexicons

```
[ ] 1 from io import StringIO
2 import pandas as pd
3 import requests
4
5
6 def get_compound(row):
7     if row['PosScore'] == 0 and row['NegScore'] == 0:
8         return 0
9     return 1 if row['PosScore'] > row['NegScore'] else -1
10
11 swn_text = requests.get('https://raw.githubusercontent.com/aesuli/SentiWordNet/master/data/SentiWordNet_3.0.0.txt')
12 columns = ['POS', 'ID', 'PosScore', 'NegScore', 'SynsetTerms', 'Gloss', 'CompoundScore']
13 df_swn = pd.read_csv(StringIO(swn_text.text), header=None, names=columns)
14
15 df_swn['CompoundScore'] = df_swn.apply(get_compound, axis=1)
```



```
[ ] 1 df_swn.loc[100:110]
```

	POS	ID	PosScore	NegScore	SynsetTerms	Gloss	CompoundScore
100	a	21403.0	0.250	0.250	unobliging	NaN	-1
101	a	21592.0	0.125	0.250	uncooperative	NaN	-1
102	a	21766.0	0.500	0.000	accurate	NaN	1
103	a	22219.0	0.625	0.000	faithful	NaN	1
104	a	22437.0	0.625	0.250	dead-on	NaN	1
105	a	22680.0	0.500	0.125	high-fidelity	NaN	1
106	a	22852.0	0.500	0.000	surgical	NaN	1
107	a	22962.0	0.000	0.000	straight	NaN	0
108	a	23120.0	0.125	0.000	true	NaN	1
109	a	23278.0	0.875	0.000	veracious	NaN	1
110	a	23383.0	0.125	0.625	inaccurate	NaN	-1

Features

- Filter by adjectives, nouns, verbs and adverbs
- Contextual valence shifters: negation, intensification
 - NOT good/NOT bad
 - "VERY good", "EXTREMELY good"
 - "I REALLY love the weather." vs. "I HARDLY love the weather."
- Phrases/N-Grams
- Semantic Expansion: WordNet
- Emoticons
- Word Embeddings

PoS-tag Filtering

```
[ ] 1 from nltk import pos_tag
    2 from nltk import word_tokenize
    3 import nltk
    4 nltk.download('punkt')
    5 nltk.download('averaged_perceptron_tagger')
    6
    7
    8 pos_tags_noun = {'NN', 'NNS', 'NNP', 'NNPS'}
    9 pos_tags_adj = {'JJ', 'JJR', 'JJS'}
```

```
[ ] 1 sentence = 'The quick brown fox jumps over the lazy dog.'
    2 tokens = word_tokenize(sentence)
    3 tagged = pos_tag(tokens)
    4 tagged[:4]
```

```
☞ [('The', 'DT'), ('quick', 'JJ'), ('brown', 'NN'), ('fox', 'NN')]
```


```
[ ] 1 [token[0] for token in tagged if token[1] in pos_tags_noun]
```

```
☞ ['brown', 'fox', 'dog']
```

```
[ ] 1 [token[0] for token in tagged if token[1] in pos_tags_adj]
```

```
☞ ['quick', 'lazy']
```

N-Grams



WordNet

```
[ ] 1 from nltk.corpus import wordnet as wn
    2 import nltk
    3
    4 nltk.download('wordnet')
```

```
[ ] 1 print(f"Synsets: {wn.synsets('sentiment')}")
    2 print(f"Definition: {wn.synset('opinion.n.01').definition()}")
    3 print(f"Hypernyms: {wn.synset('opinion.n.01').hypernyms()}")
    4 print(f"Hyponyms: {wn.synset('opinion.n.01').hyponyms()[:2]}")
```

↳ Synsets: [Synset('sentiment.n.01'), Synset('opinion.n.01')]
Definition: a personal belief or judgment that is not founded on proof or certainty
Hypernyms: [Synset('belief.n.01')]
Hyponyms: [Synset('eyes.n.01'), Synset('idea.n.03')]

```
[ ] 1 print(f"Similarity opinion<->sentiment:", wn.synset('opinion.n.01').path_similarity(wn.synset('sentiment.n.01'))))
    2 print(f"Similarity car<->truck:", wn.synset('car.n.01').path_similarity(wn.synset('truck.n.01'))))
    3 print(f"Similarity car<->bed:", wn.synset('car.n.01').path_similarity(wn.synset('cat.n.01'))))
```

↳ Similarity opinion<->sentiment: 0.1
Similarity car<->truck: 0.3333333333333333
Similarity car<->bed: 0.05555555555555555

Features - An Evaluation

- Content-based features
 - Bag of words (BOW)
 - Word embeddings (W2V)
 - Concept embeddings (C2V): map text onto concepts from the UMLS Metathesaurus
- Domain-specific features
 - Extraction UMLS Semantic Types (ST)
- Positional features (Pst)
 - position of the sentence within the post, position of the post within the thread
- Network features (Net)
 - number of replies of the post, Is a primary question (1 if the sentence belongs to the first post in a conversation, and 0 otherwise)

Features - An Evaluation

- Sentiment-based features (SA)
 - Number of positive/negative words (General Inquirer], SentiSense, and SentiStrength)
 - Emotions expressed in the text (from SentiSense lexicon: Ambiguous, Anger, Calmness, Despair, Disgust, Anticipation, Fear, Hate, Hope, Joy, Like, Love, Sadness, and Surprise)
- Grammatical features (Gramm)
 - Verb tense
 - Part-of-speech
 - Negation
- Factuality (Fact)
 - Includes the factual/non-factual label manual classification as a feature

	SMO		RF		NB		Vote	
Feature	Acc	F-1	Acc	F-1	Acc	F-1	Acc	F-1
BOW	57,1	47,7	57,6	46,1	58,1	49,4	58,4	48,9
BOW + ST	60,9	60,1***	64,4	59,4	63,9	62,9	63,4	61,5
BOW + ST + Pst	61,8	59,7***	64,8	59,9	63,6	62,6	63,9	62,4
BOW + ST + Pst + Net	59,8	59,3***	64,7	59,9	63,5	62,7	62,9	61,3
BOW + ST + Pst + Net + SA	62,1	61,6***	64,2	59	65,3	64,5	65,5	64,1
BOW + ST + Pst + Net + SA + Gramm	56,6	47,1	57,2	55,5	57,6	49,2	59,3	52,4
BOW + ST + Pst + Net + SA + Gramm + Fact	57,1	49,2	56,7	55,3	57,5	49,7	59,7	53,6
W2V	65,9	65,1	62,8	56,5	57,5	46,4	66,5	63,5
W2V + ST	66,2	65,2	63,1	56,4	60,3	62,9	67	63,4
W2V + ST + Pst	66,5	65,9	61,3	53,8	60,6	62,6	67,3	63
W2V + ST + Pst + Net	66,1	65,4	62,3	55,2	59,9	62,7	67,4	63,4
W2V + ST + Pst + Net + SA	68,4	66,5	62,6	55,7	61,9	64,5	68,7	65,3
W2V + ST + Pst + Net + SA + Gramm	66	65,8	62,9	56,5	62,6	49,2	68,2	65,9
W2V + ST + Pst + Net + SA + Gramm + Fact	64,7	64,5	63,3	56,9	62,6	49,7	68,5	66,4
C2V	62,6	60,8	61,2	55	57,4	45,6	63,8	59,6
C2V + ST	65	65,2**	63,1	56,4	63,9	53,8	65	61,4
C2V + ST + Pst	66,5	65,9**	61,3	53,8	63,7	54,4	65,3	62
C2V + ST + Pst + Net	66,1	65,4**	62,3	55,2	63,5	54,1	66,4	63,4
C2V + ST + Pst + Net + SA	67,8	67,2**	62,6	55,7	65,3	57,6	67,7	65,3
C2V + ST + Pst + Net + SA + Gramm	62,4	61,6*	63,1	56,5	57,6	58,8	64,8	62
C2V + ST + Pst + Net + SA + Gramm + Fact	63,2	62,4*	63,3	57	57,5	59,5	65,3	62,5

<https://doi.org/10.1371/journal.pone.0207996.t008>

Feature engineering for sentiment analysis in e-health forums: journals.plos.org/plosone/article?id=10.1371/journal.pone.0207996

Features with Deep Learning

```
[ ] 1 import re
    2 from nltk.stem import WordNetLemmatizer
    3 from nltk.corpus import stopwords
    4
    5 stop_words = set(stopwords.words("english"))
    6 lemmatizer = WordNetLemmatizer()
    7
    8
    9 def clean_text(text):
   10     try:
   11         text = re.sub(r'^\w\s', '', text, re.UNICODE)
   12     except:
   13         print(text)
   14     text = text.lower()
   15     text = [lemmatizer.lemmatize(token) for token in text.split(" ")]
   16     text = [lemmatizer.lemmatize(token, "v") for token in text]
   17     text = [word for word in text if not word in stop_words]
   18     text = " ".join(text)
   19     return text
```

Evaluation and Limitations

Evaluation

- Recall
 - measure of completeness: $TP/(TP+FN)$
- Precision
 - measure of correctness: $TP/(TP+FP)$
- F-measure
 - harmonic mean of recall & precision: $(2*recall*precision)/(recall+precision)$
- Accuracy
 - proportion of true results among in all cases: $(TP+TN)/(TP+TN+FP+FN)$
- Cross-validation

Evaluation

```
[ ] 1 from sklearn.metrics import classification_report
    2 y_true = [1, 1, -1, 1, -1, -1, 1]
    3 y_pred = [1, -1, -1, 1, 1, -1, 1]
    4
    5 print(classification_report(y_true, y_pred))
```



	precision	recall	f1-score	support
-1	0.67	0.67	0.67	3
1	0.75	0.75	0.75	4
accuracy			0.71	7
macro avg	0.71	0.71	0.71	7
weighted avg	0.71	0.71	0.71	7

Limits

- Context, word ambiguity
 - This movie's plot is unpredictable.
 - The brakes of this car are unpredictable.
- Stilistic means: sarcasm, irony, cynicism, metaphors
- Domain-dependency
- Colloquial language (e.g. in social media)
 - "I luv u"
 - NLTK TweetTokenizer()

Corpus Annotation

Corpus Annotation

- Choose level of granularity
- Split corpus:
 - Set for annotation guidelines
 - Set for main annotation task
- Derive guidelines on small set
- Select annotators: level of expertise required?
- Annotators annotate on large set using the guidelines

Corpus Annotation

- Calculate inter-rater agreement
 - Cohen's kappa: for 2 raters
 - Fleiss' kappa: more than 2 raters

```
[ ] 1 from sklearn.metrics import cohen_kappa_score
```

```
[ ] 1 rater_1 = [1, 1, 1, 0, 0, 1, 1, 0, 0, 1]  
  2 rater_2 = [0, 1, 1, 0, 0, 1, 1, 0, 0, 1]  
  3  
  4 cohen_kappa_score(rater_1, rater_2)
```

```
↳ 0.8
```

< 0: Agreement below chance

0.01-0.2: Small agreement

0.21-0.4: Fair agreement

0.41-0.6: Moderate agreement

0.61-0.8: Substantial agreement

0.81-0.99: Almost perfect agreement

Corpus Annotation

1. If at least two annotators assigned the same label to the sentence, then such label was finally assigned.
2. If each of the three annotators assigned a different label to a sentence, then a fourth annotator was asked to select the final label.
3. If a sentence was not labeled by at least two annotators, we preserve the sentence for readiness and labeled it as NOT_LABELLED.

Corpus Annotation

	Experience	Opinion	Fact
Allergies	86%	69%	65%
Crohn	88%	65%	72%
Breast cancer	77%	70%	79%
Average	84%	68%	72%

<https://doi.org/10.1371/journal.pone.0207996.t005>

	Positive	Negative	Neutral
Allergies	79%	76%	71%
Crohn	68%	68%	79%
Breast cancer	77%	70%	79%
Average	75%	71%	76%

<https://doi.org/10.1371/journal.pone.0207996.t004>

Resources

Resources

- Tools:

- NLTK (Natural Language ToolKit)
 - www.nltk.org
- spacy.io
- TextBlob: textblob.readthedocs.io/en/dev/
- VADER (Valence Aware Dictionary and sEntiment Reasoner);
 - github.com/cjhutto/vaderSentiment

- Hubs:

- English:
 - github.com/xiamx/awesome-sentiment-analysis
- German Sentiment Analysis:
 - sites.google.com/site/iggsahome
 - Offensive language: projects.fzai.h-da.de/iggsa/resources-tools-and-literature/

Resources

- Sentiment Lexicons:

- SentiWordNet:
 - github.com/aesuli/SentiWordNet
- Opinion Lexicon:
 - www.cs.uic.edu/~liub/FBS/sentiment-analysis.html
 - github.com/shekhargulati/sentiment-analysis-python/tree/master/opinion-lexicon-English
- SenticNet: sentic.net
- General Inquirer: www.wjh.harvard.edu/~inquirer/
- SentiStrength: sentistrength.wlv.ac.uk/
- SentiSense: nlp.uned.es/~jcalbornoz/SentiSense.html

- Stopwords:

- [github.com/stopwords-iso/](https://github.com/stopwords-iso)
- `from nltk.corpus import stopwords`

Thank you.

Emotion is a sum totaled by an adding machine of the mind.

Ayn Rand, *Atlas Shrugged*.